

Guessing the Washington Michelin stars

Bertrand Dolimier

September 28, 2016

Executive Summary

Objective: Predict the result of the newly created Michelin guide for the Washington, DC area. The outcome will include a) the number of restaurants starred, the name of restaurants with their corresponding stars (1, 2 or 3)

Approach: Collect data for Restaurants in the other US cities where a Michelin guide already exist, explore for each of them the publicly available information, select the most relevant dataset and come up with a predictive algorithm (FYI no personal opinion just data science).

Tech stack Used R and Knit Markdown (Rmd extension) with R libraries (httr, jsonlite & XML) to integrate documentation with code and generate the pdf doc.

The multi-step process is described below:

Base Geo. Dataset

1. Enter General area demographic in csv file with the following data point: Number of Restaurants in city, Number of Restaurants in the metro area, City Population, Metro population, Average Income and number of starred restaurants.
2. Run simple model to predict number of stars to be awarded in the Washington DC area.

Enriching the 2 datasets

1. Collect the names of currently awarded restaurants along with the number of stars for New-York, Chicago and San-Francisco.
2. To better predict a sample of non-awarded restaurants are randomly selected. They were labeled as 0 starred restaurants. All restaurants of existing Michelin US areas are saved in a .csv file and will be mentioned as “training” data.
3. Collect restaurants widely selected that are potential candidates for a Michelin award in the Washington DC area
4. Collect in the preeminent local press the professional gastronomy reviews. We will collect the number of review along with the year of the review and the number of time the restaurants was mention in the dailies archives. The news organism selected were the New-York Times, the Chicago Tribunes and the San-Francisco Chronicle.
5. Collect customers reviews and rating from the Yelp web site.

Polishing the DC dataset

1. Removing the bib gourmand from the test dataset.
2. Removing entries with no reviews


Analysing the Data

1. Use linear regression to determine a model based on the data collected above.
2. Collect restaurants from the Washington area.
3. Run the model selected and order the restaurants.
4. Using demographic information and number of restaurant in each collective area weight by the average customer rating estimate the number of 3,2 and 1 stars restaurant to be anticipated.
5. Give an arbitrary value to the Restaurant type based on occurrence in training dataset
6. Assign a “novelty factor” to the Washington DC area as less restaurant can be expected the first year a guide is publish
7. Ranked list and determine the “winners”...

Detail and code

Base Geo Dataset

The simple plot below, shows the correlation between starred restaurants with the selected multiple criteria that are: City Population, Metro Population, Income and Number of “Quality” restaurants. San Francisco appears at an outlier in several ways (not for this discussion)



(code hidden)

Collect other local restaurants in Existing US Michelin guide areas,
and save as training dataset

```
name = paste( c( "/Users/bdolimier/persodev/michelinDC/output/usa2016_result.csv" )  
usadata <- read.csv( name , header=TRUE, stringsAsFactors=FALSE , sep = "," )  
dfusa <- as.data.frame( usadata )  
dfusa <- dfusa[ -c(1) ]  
myIndex <- length( dfusa[[1]])  
  
# Take zz pages sample...  
for ( zz in 1:4 ) {  
  if (zz==1) zone <- "New-York"  
  if (zz==2) zone <- "Chicago"  
  if (zz==3) zone <- "San Francisco"  
  if (zz==4) zone <- "Washington"  
  zoneEncode <- URLencode( toString( zone ) )  
  
  # Get sorted by rating the 2 price points 4 and 3  
  for ( price in 3:4 ) {  
    for ( page in 0:15 ) {  
      start <- page * 10  
      YcacheName <- paste0( "/Users/bdolimier/persodev/michelinDC/cache/" , "YELP_COLLECT_" , zone,"_", price, "_", page )  
  
      # Reading the cache  
      html.raw<-htmlTreeParse( YcacheName, useInternalNodes=T )  
      nbc <- nchar( as( html.raw , "character") )  
      nbc  
      # Not cached let's ask Yelp  
      if ( nbc < 300 ) {  
        query=paste0( "https://www.yelp.com/search?find_desc=Restaurants+&find_loc=",zoneEncode,&start=page*&price="+price )  
        download.file( query , destfile = YcacheName , method="curl")  
        html.raw<-htmlTreeParse( YcacheName, useInternalNodes=T )  
      }  
  
      bizName <- xpathSApply(html.raw , "//span[@class='indexed-biz-name']", xmlValue)  
      bizName <- gsub( "Restaurant" , "" , bizName )  
      bizName <- gsub( "\n" , "" , bizName )  
      bizName <- gsub( "\\." , "" , bizName )  
      bizName <- gsub( " " , "" , bizName )  
      nb <- length( bizName )  
      if ( nb > 0 ) {  
        for ( jj in 1:nb ) {  
          bizName[jj] <- gsub( paste0(start+jj," ") , "" , bizName[jj] )  
          myIndex = myIndex + 1  
          dfusa[ myIndex ,] <- c("", "", "", 0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0)  
          dfusa$Restaurant[myIndex] <- bizName[jj]  
          dfusa$Zone[myIndex] <- zone  
          dfusa$Ypricing[myIndex] <- price  
          print ( bizName[jj] )  
        }  
      }  
    }  
  }  
}
```



```

    }
  }
}

View( dfusa )
write.csv( as.data.frame(dfusa) , file = "/Users/bdolimier/persodev/michelinDC/output/michelin2016_wash

```

Get Restaurant press review from NewYork Times, Chicago Tribunes & SanFrancisco Chronicle

Using google apis Populating the fields: dfusa\$art2016, art2015, art2014, artBefore (older than 2014), dateUNKNOWN

```

## Engine id:
dcId  <- "faxe8adyxke" # - Washington Post = 017019937838437061749:faxe8adyxke
nyId  <- "uspap9gs9vq" # - New York Times = 017019937838437061749:uspap9gs9vq
chiId <- "duuro39maos" # - Chicago Tribunes = 017019937838437061749:duuro39maos
sfId  <- "navjubg8x6w" # - San Francisco Chronicle = 017019937838437061749:navjubg8x6w

name = paste( c( "/Users/bdolimier/persodev/michelinDC/output/usa2016_training.csv" ) )
usadata <- read.csv( name , header=TRUE, stringsAsFactors=FALSE, sep = "," )
dfusa <- as.data.frame( usadata )
dfusa <- dfusa[ -c(1) ]

iimin = 1
iimax = nrow( dfusa )

for(ii in iimin:iimax ) {
  resto <- toString(dfusa$Restaurant[ii])
  restoEncode <- URLencode( toString( paste("Restaurant",resto ) ) )
  zone = toString(dfusa[ii,3])
  if ( zone == "New-York" )      engineId = nyId
  if ( zone == "Chicago" )      engineId = chiId
  if ( zone == "San Francisco" ) engineId = sfId
  print( paste( restoEncode, zone ) )
  ns <- 70 # init to large than we will query for / 7 pages

  for ( page in 0:6 ) {
    cacheName5 <- paste0( "/Users/bdolimier/persodev/michelinDC/cache/" , "C5_", zone,"_", ii , "_" ,
    cacheName6 <- paste0( "/Users/bdolimier/persodev/michelinDC/cache/" , "C6_", zone,"_", ii , "_" ,
    cf <- NULL
    cf5 <- NULL
    tryCatch( cf <- fromJSON(file( cacheName6 )) , error = function(e) cacheFlg <- FALSE , warning =
    tryCatch( cf5 <- fromJSON(file( cacheName5 )) , error = function(e) cacheFlg <- FALSE , warning =

    ## If not already cached got ask google
    if ( is.null(cf) == TRUE || is.null(cf5) == TRUE ) {
      start = (page*10)+1
      if ( start < ns ) {
        print( paste("reading a new query", ii, page) )
        query=paste0("https://www.googleapis.com/customsearch/v1?key=AIzaSyAZBzP3NEFa7Fq76ElQW5PZaIKu

```

```

res <- GET(query)

cf5 <- do.call("rbind", content(res)[5] ) # a dataframe
write( toJSON(cf5) , file=cacheName5) # write in cache
cf5 <- fromJSON(toJSON(cf5))

cf <- do.call("rbind", content(res)[6] ) # a dataframe
write( toJSON(cf) , file=cacheName6) # write in cache
cf <- fromJSON(toJSON(cf))
}
}

# If still null skip it
if ( toString(cf) != "" && ( !is.null(nrow(cf[[1]])) ) ) {
  nc <- nrow(cf[[1]])
  if ( is.null(nc) ) nc = 0
  ns <- as.numeric( cf5[[3]] )
  if ( length(ns) == 0 ) ns = 0
  dfusa$artTotal[ii] = ns
  if ( ns>10000 ) ns = 10000
  dfusa$artTotalCap[ii] = ns

  if ( nc > 0 ) {
    for ( jj in 1:nc ) {
      titre <- cf[[1]]$title[[jj]]
      # Date method 1
      datum <- cf[[1]]$pagemap$metatags[[jj]]$ptime
      # Date method 2
      if ( is.null(datum) ) datum <- cf[[1]]$pagemap$metatags[[jj]]$date
      # Date method 3
      if ( is.null(datum) ) datum <- cf[[1]]$pagemap$metatags[[jj]]$sailthru.date
      # Date giving up
      if ( is.null(datum) ) {
        print ( paste("no date for ", ii, page, jj))
        datum = -1
      }

      dd <- substr( toString( gsub("-", "", datum) ), 0, 4)
      print ( paste( ii , page , jj ,datum) )
      if ( dd == 2016 ) dfusa[ii,5] <- dfusa$art2016[ii]+1
      if ( dd == 2015 ) dfusa[ii,6] <- dfusa$art2015[ii]+1
      if ( dd == 2014 ) dfusa[ii,7] <- dfusa$art2014[ii]+1
      if ( dd < 2014 ) dfusa[ii,8] <- dfusa$artBefore[ii]+1 # older than 2014
      if ( dd == -1 ) dfusa[ii,9] <- dfusa$dateUNKNOWN[ii]+1 # date unknown
    }
  }
}
}
}

View( dfusa )
write.csv( as.data.frame(dfusa) , file = "/Users/bdolimier/persodev/michelinDC/output/michelin2016_train

```

Get Restaurant press review from the Washington Post.

Populating the fields: dfusa\$art2016, art2015, art2014, artBefore (older than 2014), dateUNKNOWN (code hidden)

Yelp review and Classification

Populating Yreview, Yrating, Ypricing (code hidden)

Get google rating and neighborhood

Populating: Zreview, Zrating, Neighborhood (code hidden)

Zagat review through google apis


Populating fields: Fscore, Dscore, Sscore, type, buzz, buzz2, Neighborhood (code hidden)

Analyse the data

Fields collected are:

1. Restaurant: Name of the Restaurant
2. Chef: Name of the Chef (not always present)
3. Zone: New-York, Chicago, SF and DC
4. Etoile: 0 to 3 Stars
5. art2016: Number of mention in select press for the zone in 2016
6. art2015: Number of mention in select press for the zone in 2015
7. art2014: Number of mention in select press for the zone in 2014
8. artBefore: Number of mention in select press for the zone before 2014
9. dateUNKNOWN: Number of mention in select press for the zone date unknown
10. artTotal: Number of mention in select press for the zone (all dates)
11. Yreview: Number of Yelp reviews
12. Yrating: Aggregated Yelp review
13. Ypricing: Yelp pricing level ("\$" to "\$\$\$\$")
14. Zrating: Zagat rating
15. Zpricing: Zagat pricing level
16. Zbuzz: Zagat buzz indicator
17. Zbuzz: Zagat buzz indicator 2
18. Neighborhood: Restaurant Neighborhood

Create Mode / Plot Linear Regression



(code hidden)

Storing results

(code hidden)